

Amazon Simple Storage Service

Object storage built to store and retrieve any amount of data from anywhere

Amazon Simple Storage Service (Amazon S3) is a scalable object storage service designed to store, retrieve, and archive any amount of data. Amazon S3 offers various storage classes and easy-to-use management features to optimize costs, organize and analyze data, and configure fine-tuned access controls to meet specific business and compliance requirements.

Concepts:

- **Buckets:** Containers for storing objects.
- **Objects:** Files, images, and videos, and any metadata that describes those files.
- **Keys:** Unique identifiers for objects within a bucket.
- **S3 Versioning:** A configuration to keep multiple variants of an object in the same bucket.
- **Access management:** Identity and Access Management policies, bucket policies, and access control lists (ACLs).
- **S3 Lifecycle:** A configuration to manage your objects and store them cost effectively throughout their lifecycle.
- **S3 Object Lock:** A configuration that prevents Amazon S3 objects from being deleted or overwritten for a fixed amount of time or indefinitely.

How It Works

The following is an image of how Amazon S3 works. Amazon S3 stores data as objects within buckets. The data can come from various sources, such as direct uploads, AWS services, third-party applications, on-premises systems, content delivery networks, streaming data, development tools, and more.



Use Cases

Backup and storage

Static website hosting

Hybrid Cloud storage

Data lake